

SVMs for Prediction of NFL Half-Time Leads using First Quarter Data

Abstract

In this project I train support vector machines (SVMs) to predict if a National Football League (NFL) team will be leading at the half or not, given the results of the first quarter. For each team, I train a SVM using the optimal hyperparameters and features for prediction. The results show that the optimal classifying model for each team is, on average, 17.6% better than flipping a coin.

Background

Gridiron football is the most watched sport in America- and with most sports, there is an element of unpredictability that keeps us glued to the television. At the end of the 1st quarter or the half of an NFL game with my friends, we typically ask each other who will lead in the following quarters and make “friendly bets” from here. I was inspired to ask the following question after losing one of these bets: “Given information about teams before the start of the second quarter, can we accurately predict if they will lead at the half?”

A question immediately following the above, should be “What should we measure?”. With so many variables in play, it’s near-to-impossible to find the correct answer. The approach I used here was to look at the team’s performance in offense, defense, and special teams.

Football teams are traditionally divided into three “platoons”: offense, defense, and special teams. The players on offense focus on driving the football down the field, earning points if they move the ball to the opposing end-zone. The goal of the defense is to prevent the other team from scoring. Special teams come in to play when a team wants to kick the ball through the opponent’s goal, or as far as they can to the other side for the opponents to get.

It makes sense that the better a team performs in these areas, the probability that this team will take the lead increases.

Feature Selection

This project uses the 2009 to 2016 play-by-play data-set compiled by the Carnegie Mellon Sports Analytics Club. If we treat each first quarter of a regular season game (ignoring pre-season and post-season) as one data point *for each team*, we get:

$$N = 16 \text{ games} \times (2016 - 2009) = 128 \text{ points per team}$$

Following the 50:25:25 rule for splitting our data-set into training, verification, and testing sets, we get $N = 64$ *points* for training. There are over 100 features *per play* in our data-set, which means we need to restrict the information we convey. An arbitrary maximum of $d = \sqrt{N} = 8$ features per set was chosen as this restriction. The four feature sets I looked at were:

1. **Offense Focused**, $d = 7$.
2. **Defense Focused**, $d = 8$.
3. **Special Teams Focused**, $d = 8$.
4. **Overall Performance**, $d = 8$.

It is important to note that the selection of these feature ignores the other team's performance (directly), the changes made over the years, weather of this game, and much more. These factors go beyond the scope of this project, but are important to be aware of.

Methods

This is a binary classification problem, whose labels $[\ell_1, \ell_2]$ represent $\ell_1 =$ leading at the half and $\ell_2 =$ not leading at the half. To solve this classification problem, I used SKLearn's support vector machines.

Considering that each team is different, I wanted to tailor my models to each team. There are 32 teams in the NFL, which means our goal is to find 32 different models. There are three parameters I am testing: the features, the c SVM parameter, and the γ SVM parameter. The radial basis function kernel was used for all SVMs.

- Features: Offense Focused, Defense Focused, Special Teams Focused, Overall Performance
- c SVM Parameter: 0.5, 1, 10, 100, 1000
- γ SVM Parameter: 0.0001, 0.001, 0.01, 0.1, 1

To train each model, I used all games whose week number was in the space $[1, 3, 5, 7, 9, 11, 13, 15]$ (the odd weeks). This gives us 50% of our entire data-set ($N = 64$ per team) to train with, and avoids any patterns that may emerge by only choosing the early or late weeks.

After training our array of models, we need to determine which model is the best for each team. For all models belonging to a team, we predict the lead of games that don't belong to our training data-set. This new data-set consists of games whose week number was 2, 6, 10, or 14, giving us $N = 32$ points per team to verify with. The best model with the most optimal hyperparameters c and γ , as well as the optimal feature set were found using a grid search (iterating through each combination of c , γ , and our feature set). This model would then be chosen for our testing phase.

To test how accurate our final models are for each team, we predict the lead of games that don't belong to our previous data-sets. This new data-set consists of games

whose week number was 4, 8, 12, or 16, giving us $N = 32$ points per team to test with. The results of these trials are given below.

Results

Verification Data Summary

- c Parameter
 - a. 25% of best models use $c = 0.5$.
 - b. 6% of best models use $c = 1$.
 - c. 31% of best models use $c = 10$.
 - d. 10% of best models use $c = 100$.
 - e. 28% of best models use $c = 1000$.
- γ Parameter
 - a. 5% of best models use $\gamma = 0.0001$.
 - b. 13% of best models use $\gamma = 0.001$.
 - c. 3% of best models use $\gamma = 0.01$.
 - d. 13% of best models use $\gamma = 0.1$.
 - e. 66% of best models use $\gamma = 1$.
- Feature Set
 - a. 18% of best models use offense focused features.
 - b. 13% of best models use defense focused features.
 - c. 16% of best models use special team's focused features.
 - d. 53% of best models use the overall performance features.

Testing

- On average, each model was $(0.676 \pm 0.097) \times 100\%$ accurate.
- The most accurate model belongs to the Tampa Bay Buccaneers (TB), who used $c = 1$, $\gamma = 0.1$, and the overall performance feature set.
- The least accurate model belongs to the Buffalo Bills (BUF), who used $c = 1$, $\gamma = 1$, and the defense feature set.

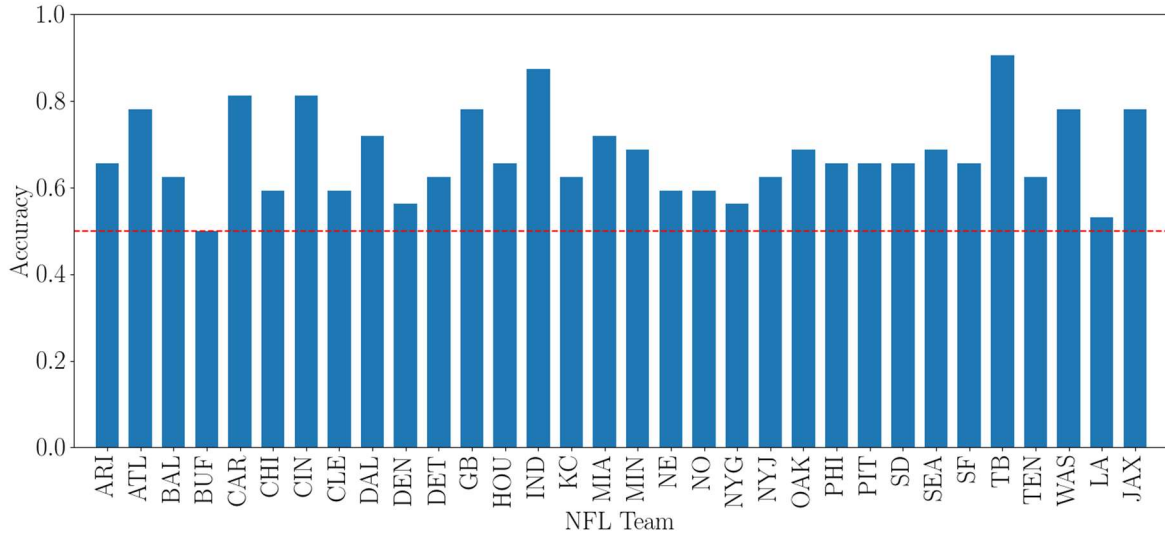


Figure 1, depicts results of testing the best classifying model for each team in the NFL.

Discussion

Most of the best models use $c = 1000$ and $\gamma = 1$, the highest of each parameter we tested. We might be able to increase our model accuracy by looking at even higher values for c and γ here. It is interesting to note that over half of our models use the overall performance feature set. This shows how every “platoon” in football matters, and how we need to look at all aspects of the game to answer our original question.

For a more involved project, I would have considered more data-sets. It’s a common assumption that teams play different in weather that is different from their home city, so I could include weather features. Maybe the people in attendance is a factor as well, and we could include the number of tickets sold as another feature. The week number itself could also be a feature as teams will play differently if they are fighting for a playoff spot, as opposed to the teams who have been eliminated from playoff contention.

Overall, most of these models are okay. These aren’t the holy grail of determining half time leads, but most of them are better than flipping a coin and guessing. That shows that the selected features for each team are worth considering, but there’s work to be done in testing even more feature sets.